

Working with large databases (RDBMS)

Jay Kelkar

Jay Kelkar

@jkelkar

Independent consultant

15 years as SAP Basis Consultant

25 years working with databases

Requirements

What we want from a database system:

1. Accuracy - consistent data
2. Safety - data stays available
3. Performance - for work + reporting

A modern RDBMS

Provides ACID compliance

In most cases, highly reliable

Works out of the box, for small instances

Larger systems need to be managed

A C I D

Atomicity

All changes are done or none!

No partial changes

Transactions - across multiple data

A C I D

Consistency

Satisfies all constraints

Triggers

Cascade - deletes/updates

A C I D

Isolation

Parallel or serial operations has same outcome

Assuming all changes succeed

A C I D

Durability

Data will continue to exist after a
completed transaction
- even after a shutdown

Priority

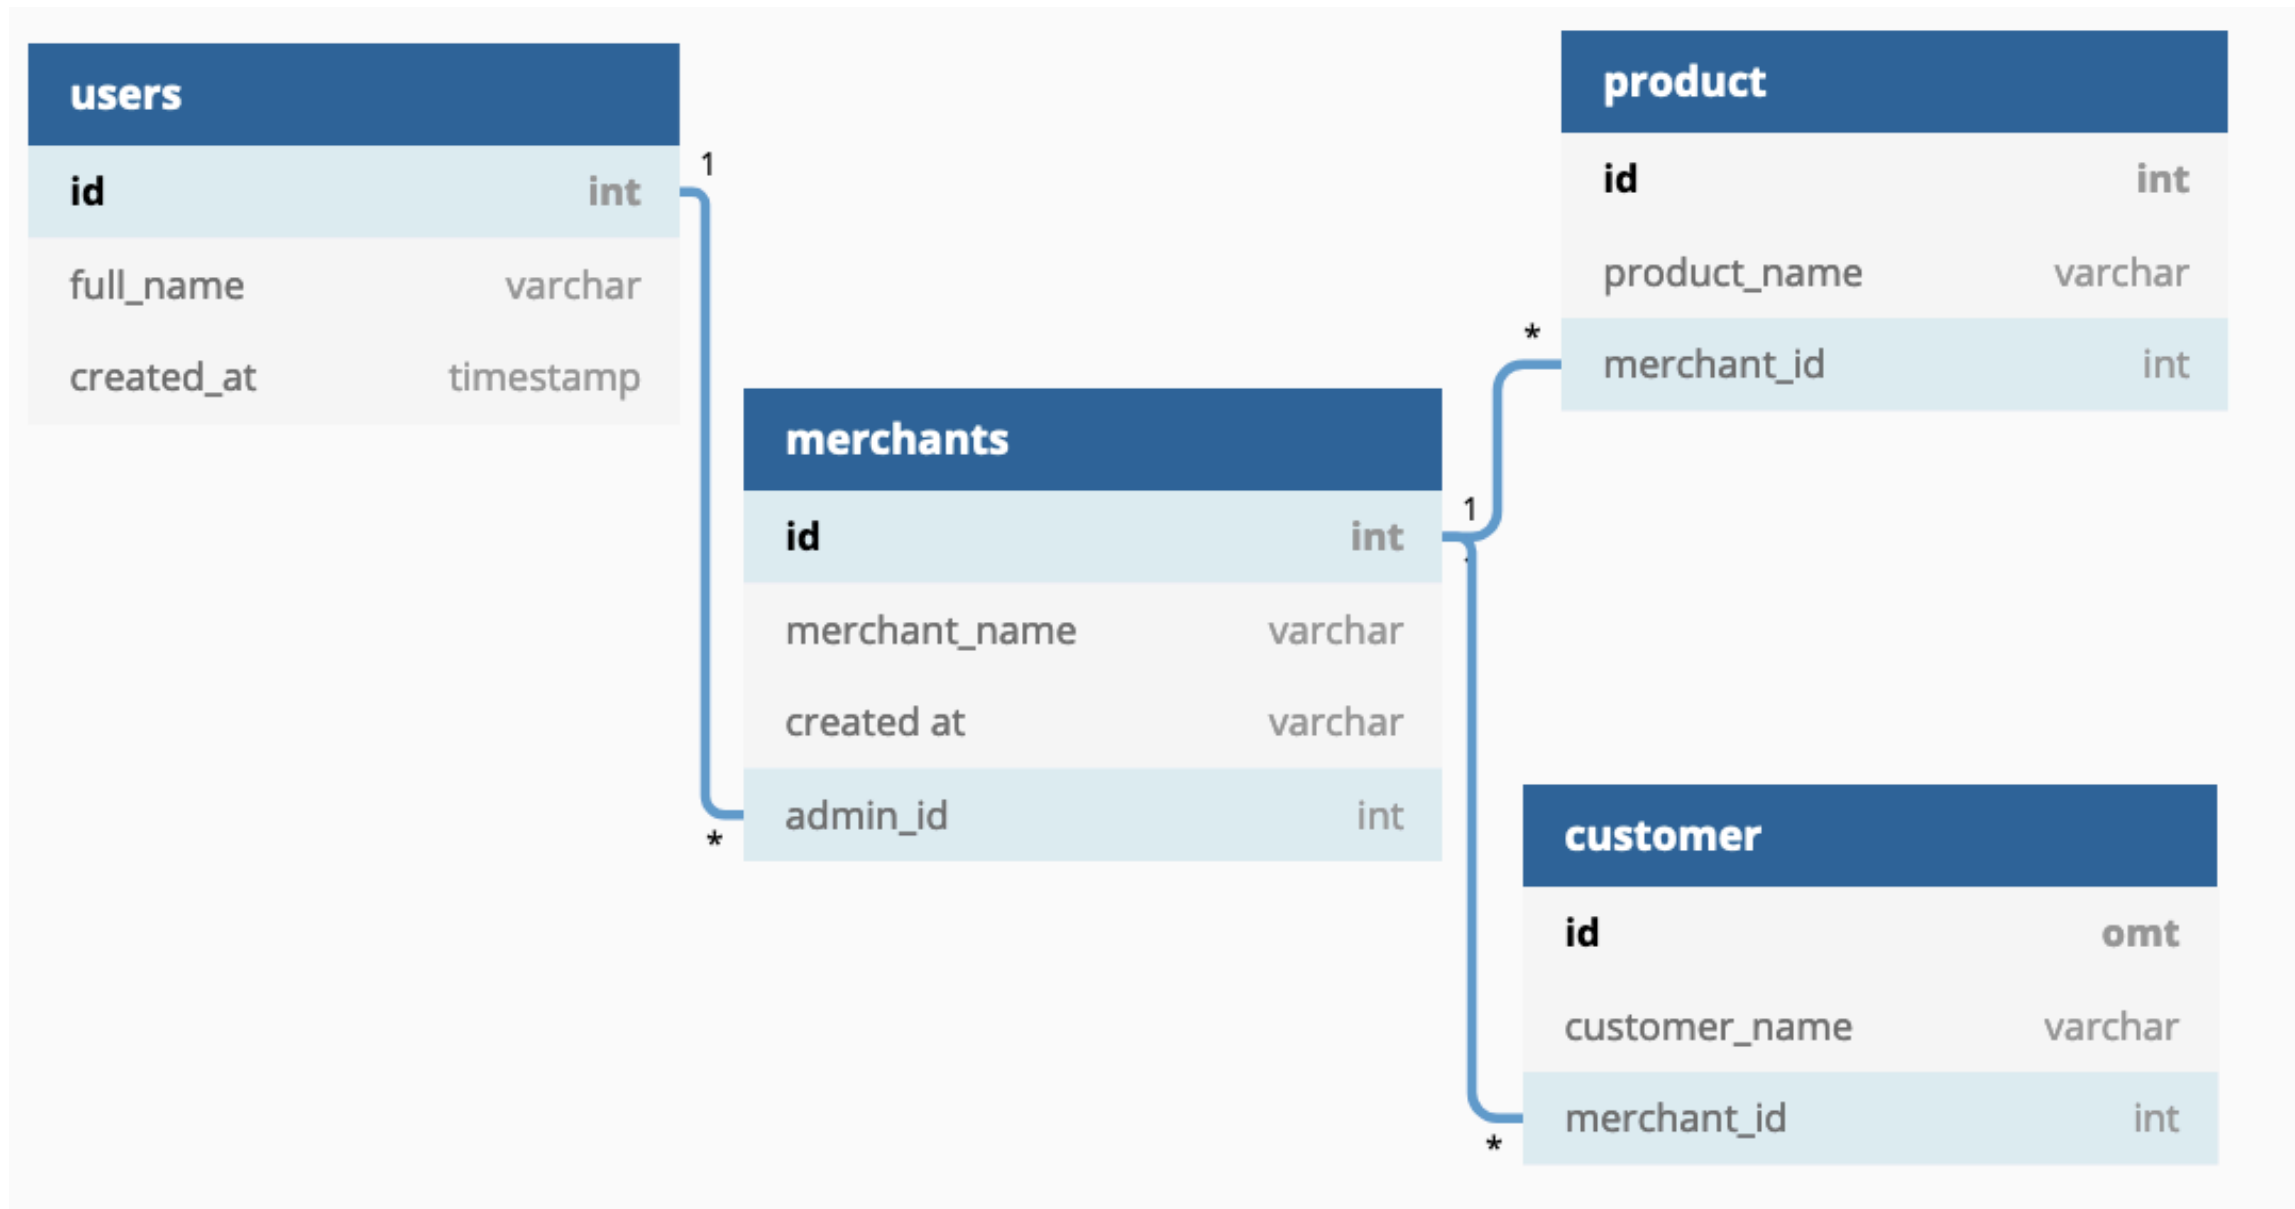
The order of requirements has to be:

- Accuracy and safety - non negotiable
- Performance - important, not primary

Accuracy

Relationships

- Foreign key relationships
- Triggers
- Cascade deletes/updates



Relationships

Accuracy

Relationships

Foreign key relationships

Field data in foreign key columns exists
only if it exists in the master table

Uses:

- master table access
- cascade deletes

Accuracy

Relationships

Triggers

Insert/update data into table following
some action

- Relieves programmer of default tasks
- Audit tables
- Keeps data consistent

Accuracy

Relationships

Cascade deleted/updates

Automatically delete or update rows
based on Foreign Key relationships

Uses:

- (delete) prevents orphan rows
- (update) set values in tables with FK on updates - in orphaned row

Safety

Backup and Restore

- Full backups
- Hot backups
- Redo log backups

Safety

Full backups

- At reasonable intervals - weekly?
- Shut down database & copy all database files
- Lock database and dump structure & data as SQL (only for smaller databases)

Safety

Full backups

Pros:

- File backups are easiest to restore
- SQL backups essentially create all DB objects and load data (insert/copy data)

Cons:

- File backups require system to be quiescent
- SQL backups require database to be locked

Safety

Hot backups

Pros:

- Can be run when DB is productive

Cons:

- Can take longer
- Must put database in backup mode
- By itself cannot restore a database

Safety

Redo log backups

Pros:

- Only way to do a point in time recovery
- Fully supported by DB software

Cons:

- Requires careful sizing of redo log files
- Need to backup redo log files often

Safety

Saving backups

- * Must keep remote copies of backup data
- * Ideally, have backup data pulled off the system. Adds a level of safety
- Ideally keep 2 copies of redo logs on remote

Safety

Durability

- Redo log file sized on risk level (mins.)
- File size is $\sim 1.5 - 2X$ for size of redo entries for a given time interval for busy system

Copy off a redo log file to backup ASAP

- smaller file size means more files
- * Force log switch before hot backup

*Verify Backup integrity

- Must verify integrity of backup
- Only way to know if a backup is useable!
- Check regularly - Monthly?

Safety

*Limit Access

- Limit # users with direct access to DB
- ALL other access through known/
tested code - **no** exceptions

Safety

Changing table structure

Suggestion:

- Prefer not to "ALTER TABLE" directly

Instead:

- Dump, drop, create and restore

A lot higher safety + a copy of data

Safety is highest priority

“Happiness is hard won, but unhappiness comes easily”

–*Anonymous*

Performance

Query execution

- A query is analyzed
- Execution cost is estimated - the plan
- The plan with the lowest cost is chosen

* Explain queries

Run "explain" on queries

Pros:

- Shows expected I/O from DB
- Shows if/which indexes are being used

Cons:

- Have to understand output
- Results can change over time

Performance

* Create needed indexes

Why do we need indexes?

- Access a data row - unique key index
- Get to part of data - sorted index
- Complex query support - multi column
- Avoid full table scan

Performance

Index

- By default PRIMARY INDEX exists - on the unique KEY column
- An index can be non unique
- Only create an index for a use case
- Adds cost to inserts, updates & deletes
- Multiple columns can be in one index

Performance

More about indexes

- Column order matters
- Query and Index columns must match

Option:

- Adjust Index to suit query
- Create a new index


Performance

Maintain indexes


Indexes often are kept as b-trees

Over time they can get unbalanced

Balanced



Unbalanced



* Maintain indexes

Maintenance:

- Keep indexes up to date - drop/create or reindex [database | table | index]
- Check if index is relevant - cost Vs. benefit

Performance

Index planning

For high performance - adjust fill factor
- space used in each leaf page



Affects indexing or reindexing

Empty space used for insert/updates

Statistics

Metadata about table data

- # unique entries per table column

Gotcha:

- Does not update automatically
- Out of date statistics result in bad query performance

* Update Statistics

Command varies by database:

- Typically you ANALYZE a table
- May use "ESTIMATE" or "COMPUTE"
- Analyzing 10-20% of data in a huge table is sufficient

* Log slow queries

A must with customer devs writing queries

- Queries can get slow over time
- Due to out of date statistics
- Due to query/index changes

Performance

Use parametric queries

If possible:

- "prepare" statements
- Use "stored procedures" & functions for Higher performance
- Can use existing execution plan

Performance

Good design

- * Balance between "normalization" and "joining" tables
- Plan "how much" data has to be read for queries
- * Use "views" when possible

Performance

DB Tuning

DB Buffers - various names

- about 25% of memory if available
- Larger amount if dedicated db server

Using a tuning specialist can give high returns in improved performance

Data access speed

Keep data on the fastest device

- * Have separate caches on App Server
- * Secondary caches restricted to

"most used data", etc.

e.g. Redis for certain tables

Performance

Table fill factor

For high performance - adjust table
fill factor

- only controls "insert" fill, rest for
updates

lower - 60+

higher ~ 100

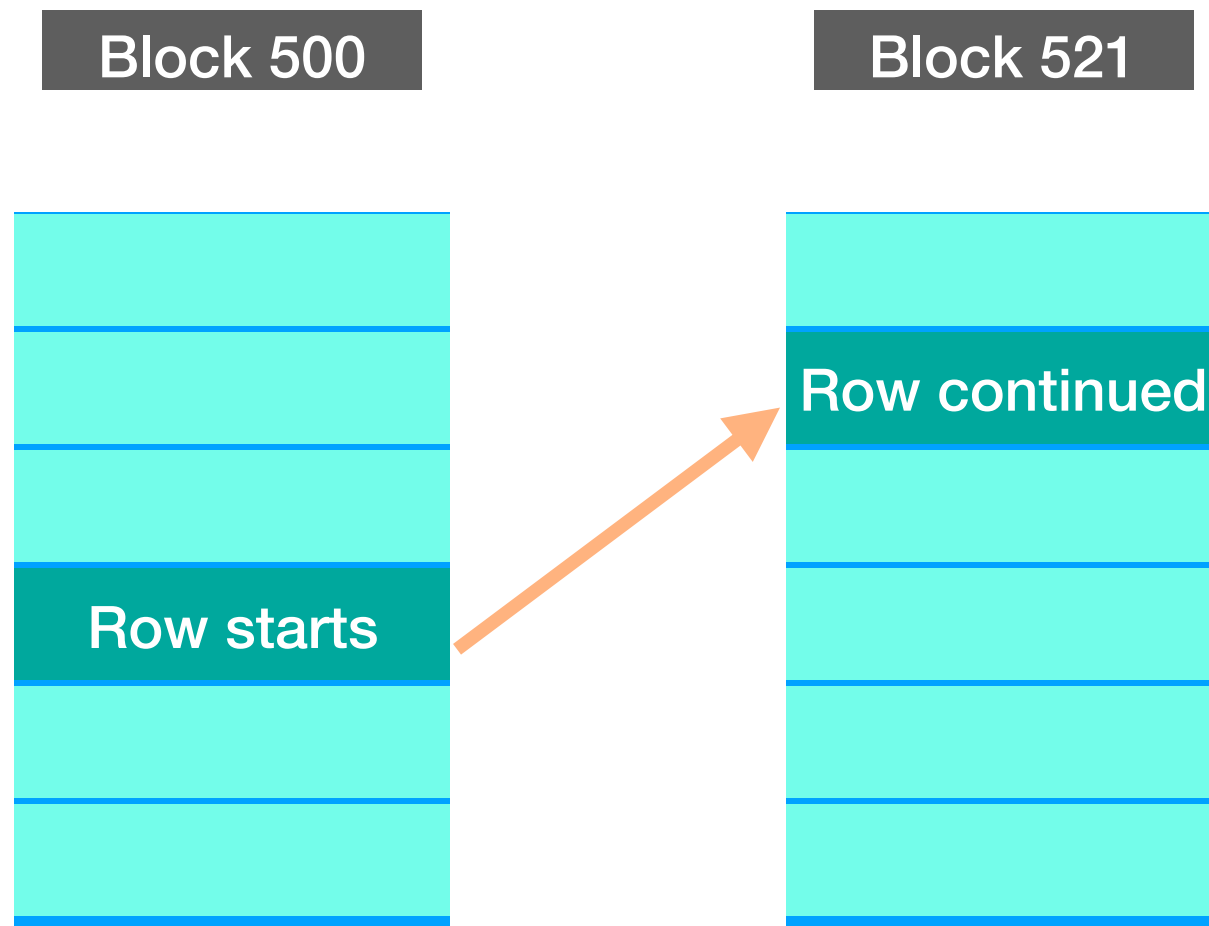
Dynamic
Tables



Static
Tables

Performance

- prevents "row chaining"



Types of data stored

Transactional data:

- most data - OLTP

Non transactional data (need not log):

- derived and repeated data
- analytics, metrics, statistics, OLAP
- anything that can be recreated

Track Performance

- slow queries
- tracking backup: full, hot and redo
- Cache hits - should be high > 90-95%

Can tell health of the system

- figure out the working set size for your application (and set buffer size)

Performance

* Offload DB disk I/O

- Offload DB disk access as possible
- Implement large caches (db buffers)
- Implement application caches

Have to deal with cache invalidations

Performance

Selected data

- * Select no more data than needed

Unnecessary data:

- Uses more DB cache, less efficient (LRU)
- Increases I/O and CPU cost of processing
- Increases data processing in application

Some suggestions

- * Use cryptographically secure passwords (size matters)
- Use secure connections between machines
- Limit select set size (possible?)
- Could prevent large amount of data loss during an attack

Performance

Reporting

- If possible report on aggregate data
- Disk space is cheap - time feels expensive
- Cache quality is not sacrificed
- Aggregate data during slow times

Performance

DB I/O Patterns

- Isolate Table data I/O from Index I/O
- Typically separate tablespaces (areas)

Keep separate network data flow
for DB data and User data

General

Version control

- Consider versioning your data objects, esp. Tables, Views and Indexes
- Makes it easy to keep track of changes
- Ideally set up a data dictionary to manage all db objects

Overall

Relative importance

- In everything we have seen, aim for high accuracy and high safety
- In the worst case, we can fudge on performance
- This gives us a reliable system, and
- Performance can always be improved

A person is riding a horse across a vast, open field at sunset. The sun is low on the horizon, creating a warm, golden glow across the sky and the ground. The sky is filled with soft, wispy clouds. The rider is silhouetted against the bright light of the setting sun. The horse is galloping, and a cloud of dust is kicked up behind its legs. In the distance, a simple wooden bench is visible on the right side of the field.

The End